



CLEARANCE

Gate the action, not the answer.

The decision layer above AI detection — deciding on the reversibility of the action, not the content of the text.

Business Proposal

Accenture Innovation Challenge 2026 · Round 2 · Track 1 (ControlPlane.ai)

Guardrails detect. They don't decide.

Today's guardrails answer “is this text risky?” — scoped by endpoint at best. But the same risky sentence is harmless in a brainstorm and dangerous when it moves money. Detection alone cannot tell the difference.

1

One verdict, every context

A hallucinated refund window is treated identically whether it is shown to a user or used to issue a refund.

2

Coverage traded for cost

Vendors sample traffic to control spend — so the risky request that slips through is often the one never checked.

3

The operating point is invisible

The thresholds that govern everything are an undocumented engineering default, not an owned, auditable decision.

THE STAKES IN 2026

Agents now act, not just answer.

Refunds. Emails. Record writes. Trades. When an LLM can take an irreversible action, “the text looked fine” is not a safety story.

**The gap isn't better
detection.
It's the decision.**

Content risk ≠ action risk.

One model response. Identical detection. Two verdicts — because the reversibility of the action differs.

ONE MODEL RESPONSE

“You are entitled to a full refund of €200 for the delay, and it will be credited to your card.”

CONTEXT A

ANNOTATE

action **ticket.note**

reversible

Goes into a draft note. Flagged with a caveat and allowed through.

CONTEXT B

ESCALATE

action **refund.issue**

irreversible

About to move money. Blocked and routed to a human for review.

Nobody else ships the paired verdict. It is the first thing our demo shows.

A decision layer that sits above detection

CLEARANCE is a drop-in, OpenAI-compatible gateway. Point your app's base URL at it and every response is governed — text streams instantly, only side effects wait at the gate.

1

Reversibility-indexed verdicts

We decide on (risk, reversibility of the action), not content alone. Identical text, two verdicts — annotate a draft, block a refund.

2

Cost-proportional coverage

100% coverage at a no-LLM tier; spend escalates only when uncertain. The expensive tier trains the cheap one, so paid traffic falls week over week.

3

Signed operating point

The thresholds are a governance artifact — chosen by a named owner, dated, and hash-stamped into every immutable ledger row.

Every detector sits behind one adapter interface — so Lakera, Patronus or Bedrock Guardrails can be the backend. We don't out-detect them; we decide above them.

One request, one decision path



Text is never held. Response tokens stream at zero added latency; only side effects wait at the latch. On a checker timeout the policy fail-mode decides — irreversible actions fail closed, reversible ones fail open.

Three things nobody else ships

1

Reversibility-indexed verdicts

Incumbents decide on content risk, scoped by endpoint. We decide on the action class — so the same sentence annotates as a draft and blocks as a refund.

*Identical text →
two verdicts, live*

2

Coverage that pays for itself

100% coverage at L0 with no model call; spend escalates only in the uncertainty band. Human labels retrain L1, so the paid tier shrinks over time.

*Paid tier 9% → 1%
over four weeks*

3

The operating point is a signed artifact

The thresholds that govern everything are owned, dated and rationalised by a named risk owner — and SHA-256-hashed into every immutable ledger row.

*Tamper one row,
the chain breaks*

We don't compete on detection — we use it

The 2026 guardrail market is crowded and good. CLEARANCE treats every incumbent as a potential backend behind its adapter interface, and adds the layer none of them ship.

PLAYER	STRENGTH	TO CLEARANCE
Lakera · OpenAI Guardrails	Input rails, injection, sub-50ms	Backend candidate
Patronus · Galileo · Bedrock	Groundedness & hallucination scoring	Backend candidate
NeMo Guardrails	Programmable dialogue / state control	Complementary
Bifrost · Agent Command Center	OpenAI-compatible gateway, policy-as-config	Closest — still content-scoped

THE ADAPTER SEAM

Detection is a commodity with excellent vendors. Our contribution is the decision layer, so every detector is swappable in one line of policy:

```
groundedness:  
  backend: patronus
```

One line, and the tiers, gate, latch and ledger are untouched. “Why not just use Lakera?” — we happily do.

A working prototype, not a slideware promise

Runs fully offline with zero API keys on a committed corpus. Clone it, run one command, and the whole system works.

1.5 ms

added p95 latency at L0, 100% coverage

6 / 6

required behaviour tests passing

9% → 1%

paid-tier traffic over four weeks, recall held

100%

offline & reproducible — zero keys

WHAT'S BUILT

- OpenAI-compatible gateway + hash-chained ledger
- 5 detectors behind a swappable adapter interface + the confidence-evidence gap
- Policy engine, action gate & latch — the live paired verdict
- Conversation accumulator — blocks an agentic action on an earlier unsupported premise
- Tuning console: operating point recomputes vs the corpus in < 300 ms

PRECISION AT THE DEFAULT OPERATING POINT

Support assistant



Internal copilot



Decision support




Precision 1.00 · recall ~0.90 · FP rate 0.00

A new budget line: AI action governance

As enterprises move LLMs from chat to agents that take real actions, spend shifts from “content moderation” to “action governance” — a control-plane budget that scales with every governed decision.

TAM	\$18B+	AI trust, safety & governance tooling by 2030 (illustrative, top-down)
SAM	\$3–4B	Runtime guardrails & policy gateways for production LLM apps
SOM	\$40–60M	Bottoms-up 3-yr wedge (see assumptions →)
Bottoms-up: ~6,000 enterprises running agentic LLMs in production × ~30% needing action-level governance × ~\$30K ACV ≈ \$54M reachable. Figures are our estimates with stated assumptions, not third-party market data.		



Priced on governed decisions, aligned with value

USAGE

Per 1,000 governed decisions

- Metered on decisions, not seats
- Cost-proportional: cheap L0 on all traffic, spend only on the uncertain slice
- Paid tier shrinks as L1 learns

PLATFORM

Annual per environment

- Tuning console + signed operating points
- Immutable audit ledger & exports
- Adapter integrations (Lakera / Patronus / Bedrock)

ENTERPRISE

Custom

- Jurisdiction packs & data residency
- Human-review workflows & SSO
- SLAs, on-prem / VPC deployment

Land & expand. Start as the governance layer on one high-blast-radius workflow (refunds, outbound email), prove the audit trail, then expand across every agent that can take an action. Net revenue retention rises as customers add governed surfaces.

Land where an action is irreversible




Fintech support

Refunds, credits, chargebacks — irreversible money movement with heavy audit and dispute exposure.



Health decision-support

Advisory that must not out-run its evidence; the lowest-appetite operating point and human escalation.



Enterprise copilots

Outbound email, record writes and HR data where a fabricated named-person detail is the signature failure.

THE MOTION



Design partners

2–3 partners on one high-stakes workflow each; free pilot, we instrument the win.



Prove the audit


Signed operating point + immutable ledger becomes the artifact their risk team wants.



Expand by surface

Every new agent action is a new governed surface — usage and NRR compound.

From prototype to control plane





CLEARANCE

Gate the action, not the answer.

Design partners

One high-stakes, action-taking workflow to govern — we run the pilot.

A reference audit

Prove the signed operating point + immutable ledger with your risk team.

The wedge

Become the action-governance layer as your agents move to production.

Prototype, docs & demo: github.com/Aparna-Jha-05/clearance

Accenture Innovation Challenge 2026 · Round 2 · Track 1

CLEARANCE